

## Final Project Code

by Group 3 - Bo Ching Liu, Soo Lee, Guanghao Li, Heng Li, Mengting Tang, Cheney Wong

```
In [1]: ▶ import pandas as pd
        from PIL import Image
        import numpy as np
        import matplotlib.pyplot as plt
        from eye_functions import *
        import plotly_express as px
        from sklearn.model_selection import train_test_split
        from skimage.measure import block_reduce
        from skimage.io import imread

In [2]: ▶ import tensorflow as tf
        from tensorflow.keras.models import Sequential
        from tensorflow.keras.layers import Dense, Flatten, BatchNormalization,
        from tensorflow.keras import regularizers
        from tensorflow.keras.callbacks import Callback
        from tensorflow.keras.preprocessing import image
        from tensorflow.keras.layers import Dense, BatchNormalization, Dropout

In [3]: ▶ eyes = pd.read_csv('full_df.csv')
        eyes = eyes[(eyes['labels']!="['N']") | (eyes['labels']!="['C']")]
        eyes['eyes_emoji']=0
        eyes.loc[eyes['labels'] == "['C']", 'eyes_emoji'] = 1

In [4]: ▶ Y = eyes['eyes_emoji']
        X = eyes['filename']

In [5]: ▶ plotter = pd.DataFrame({'Cataracts' : Y.values.astype(str) , 'eyes' : X.
        plotter.eyes.replace([1,0],[ 'Right', 'Left'],inplace=True)
        px.histogram(plotter, x='Cataracts',color='eyes',text_auto=True)

In [6]: ▶ def image_matrix(sets):
        matrices = []
        for link in sets:
            image = Image.open('preprocessed_images/%s' % link)
            if image.mode != "RGB":
                image = image.convert("RGB")
            orig_array = np.asarray(image)
            resized = block_reduce(orig_array, (2, 2, 1), np.max)
            normed_array = (resized - resized.mean())/resized.std()

            matrices.append(resized)

        stacked_arrays = np.stack(matrices, axis=0)
        return stacked_arrays
```

#

```
In [7]: #selecing only the right eye. I had hoped that perhaps the orientation of
# right_x = diabetes[diabetes['eyes_side'] ==1].filename
# right_y = diabetes[diabetes['eyes_side'] ==1].eyes_emoji
X1, x_test, Y1, y_test = train_test_split(X,Y,test_size=0.2,train_size=0.8)
```

```
In [8]: image_train = image_matrix(X1)
```

```
In [9]: image_test = image_matrix(x_test)
```


```
In [10]: print("train X shape: ", X1.shape) # X_train: numpy array with shape: (2532, 256, 256, 3)
print("train y shape: ", Y1.shape) # y_train: numpy array with shape: (2532,)
#print("test X shape: ", x_test.shape) # X_test: numpy array with shape: (512, 256, 256, 3)
#print("test y shape: ", y_test.shape) # y_test: numpy array with shape: (512,)
print("image_train.shape: ", image_train.shape)
```

```
train X shape: (2532,)
train y shape: (2532,)
image_train.shape: (2532, 256, 256, 3)
```

## Best NN model

```
In [11]: model1 = Sequential([
    Flatten(input_shape=(image_train.shape[1:])),# BatchNormalization(),
    Dense(256, activation='relu',kernel_initializer=tf.keras.initializer
    BatchNormalization(momentum = 0.2 , epsilon= 0.001),
    Dense(256, activation='relu',kernel_initializer=tf.keras.initializer
    BatchNormalization(momentum = 0.2 , epsilon= 0.001),
    Dense(256, activation='relu',kernel_initializer=tf.keras.initializer
    BatchNormalization(momentum = 0.2 , epsilon= 0.001),
    Dense(256, activation='relu',kernel_initializer=tf.keras.initializer
    BatchNormalization(momentum = 0.2 , epsilon= 0.001),
    Dense(1, activation='sigmoid')
])


model1.compile( optimizer='adam', loss='binary_crossentropy', metrics=['
```

In [12]:  model1.summary()

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 196608)	0
dense (Dense)	(None, 256)	50331904
batch_normalization (Batch Normalization)	(None, 256)	1024
dense_1 (Dense)	(None, 256)	65792
batch_normalization_1 (Batch Normalization)	(None, 256)	1024
dense_2 (Dense)	(None, 256)	65792
batch_normalization_2 (Batch Normalization)	(None, 256)	1024
dense_3 (Dense)	(None, 256)	65792
batch_normalization_3 (Batch Normalization)	(None, 256)	1024
dense_4 (Dense)	(None, 1)	257
=====		
Total params: 50,533,633		
Trainable params: 50,531,585		
Non-trainable params: 2,048		
=====		

In [13]:  model1.compile( optimizer='adam', loss='binary\_crossentropy', metrics=['

In [14]:  checkpoint\_filepath = '~\checkpoints'  
model\_checkpoint\_callback = tf.keras.callbacks.ModelCheckpoint(  
    filepath=checkpoint\_filepath,  
    save\_weights\_only=True,  
    monitor='loss',  
    mode='min',  
    save\_best\_only=True)

In [15]: `history1 = model1.fit(image_train, Y1, epochs = 72, batch_size = 64, ver`


```
Epoch 1/72
32/32 - 15s - loss: 0.6558 - binary_accuracy: 0.6612 - recall_1: 0.512
8 - val_loss: 199.2352 - val_binary_accuracy: 0.7929 - val_recall_1:
0.5128 - 15s/epoch - 468ms/step
Epoch 2/72
32/32 - 13s - loss: 0.4363 - binary_accuracy: 0.8504 - recall_1: 0.338
5 - val_loss: 21.2846 - val_binary_accuracy: 0.8363 - val_recall_1: 0.
5128 - 13s/epoch - 419ms/step
Epoch 3/72
32/32 - 13s - loss: 0.3329 - binary_accuracy: 0.8968 - recall_1: 0.276
9 - val_loss: 171.0624 - val_binary_accuracy: 0.8994 - val_recall_1:
0.2821 - 13s/epoch - 409ms/step
Epoch 4/72
32/32 - 13s - loss: 0.2878 - binary_accuracy: 0.9057 - recall_1: 0.276
9 - val_loss: 22.2475 - val_binary_accuracy: 0.8521 - val_recall_1: 0.
4872 - 13s/epoch - 405ms/step
Epoch 5/72
32/32 - 13s - loss: 0.2729 - binary_accuracy: 0.9052 - recall_1: 0.225
6 - val_loss: 65.8326 - val_binary_accuracy: 0.8777 - val_recall_1: 0.
4250 - 13s/epoch - 410ms/step
```

In [16]: `print(history1.history.keys())`


```
dict_keys(['loss', 'binary_accuracy', 'recall_1', 'val_loss', 'val_bin
ary_accuracy', 'val_recall_1'])
```

In [17]: `# Plot the training and validation loss`

```
plt.plot(history1.history['loss'])
plt.plot(history1.history['val_loss'])
plt.title('Loss vs. epochs')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')
plt.show()
```

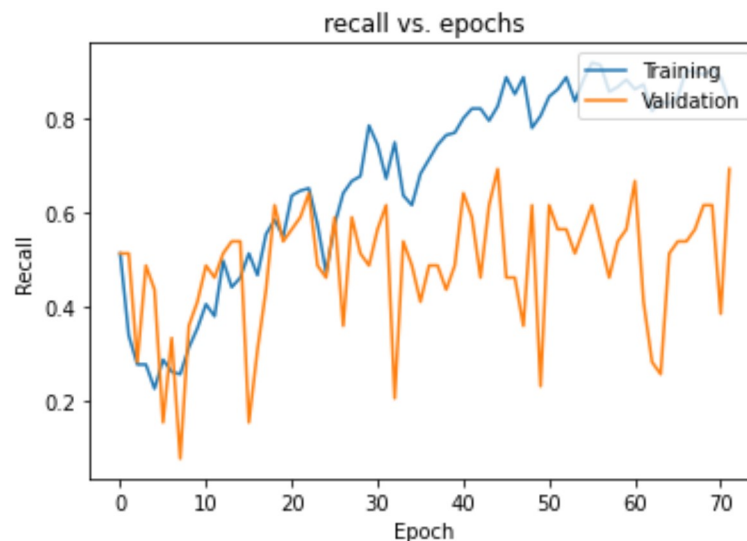


```
In [18]: ▶ plt.plot(history1.history['binary_accuracy'])
plt.plot(history1.history['val_binary_accuracy'])
plt.title('Accuracy vs. epochs')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')
plt.show()
```



```
In [19]: ▶ # Make a plot for the recall


plt.plot(list(history1.history.values())[2])
plt.plot(list(history1.history.values())[5])
plt.title('recall vs. epochs')
plt.ylabel('Recall')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')
plt.show()
```



# CNN

```
In [20]: ▶ Y1 = Y1.values.reshape((-1,1))
```

```
In [21]: ▶ model = Sequential([
    Conv2D(filters = 5, kernel_size = (4,4), activation = 'relu',
    MaxPooling2D(pool_size=2, strides=(2,2)), #2
    Conv2D(filters = 6, kernel_size = (4,4), activation = 'relu',
    MaxPooling2D(pool_size=2, strides=(2,2)), #4
    Flatten(), #5
    Dense(256, activation='relu',kernel_initializer=tf.keras.initializers.glorot_uniform),
    BatchNormalization(momentum = 0.2 , epsilon=0.001), # <- Batch Normalization
    Dense(256, activation='relu',kernel_initializer=tf.keras.initializers.glorot_uniform),
    BatchNormalization(momentum = 0.2 , epsilon=0.001),
    Dense(256, activation='relu',kernel_initializer=tf.keras.initializers.glorot_uniform),
    BatchNormalization(momentum = 0.2 , epsilon=0.001),
    Dense(256, activation='relu',kernel_initializer=tf.keras.initializers.glorot_uniform),
    BatchNormalization(momentum = 0.2 , epsilon=0.001), # <- Batch Normalization
    Dense(1, activation = 'sigmoid') #output
])
```

In [22]:  model.summary()

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 253, 253, 5)	245
max_pooling2d (MaxPooling2D)	(None, 126, 126, 5)	0
conv2d_1 (Conv2D)	(None, 123, 123, 6)	486
max_pooling2d_1 (MaxPooling2D)	(None, 61, 61, 6)	0
flatten_1 (Flatten)	(None, 22326)	0
dense_5 (Dense)	(None, 256)	5715712
batch_normalization_4 (Batch Normalization)	(None, 256)	1024
dense_6 (Dense)	(None, 256)	65792
batch_normalization_5 (Batch Normalization)	(None, 256)	1024
dense_7 (Dense)	(None, 256)	65792
batch_normalization_6 (Batch Normalization)	(None, 256)	1024
dense_8 (Dense)	(None, 256)	65792
batch_normalization_7 (Batch Normalization)	(None, 256)	1024
dense_9 (Dense)	(None, 1)	257
=====		
Total params: 5,918,172		
Trainable params: 5,916,124		
Non-trainable params: 2,048		
=====		

In [23]:  model.compile(optimizer='adam', loss='binary\_crossentropy', metrics=['bi

```
In [24]: ► checkpoint_filepath = '~/checkpoints'
model_checkpoint_callback = tf.keras.callbacks.ModelCheckpoint(
    filepath=checkpoint_filepath,
    save_weights_only=True,
    monitor='val_loss',
    mode='min',
    save_best_only=True)
```


```
In [25]: ► history = model.fit(image_train, Y1, epochs = 72, batch_size = 64, verbose=0)
```

```
Epoch 1/72
32/32 - 13s - loss: 0.6368 - binary_accuracy: 0.6706 - recall_2: 0.564
1 - val_loss: 0.6663 - val_binary_accuracy: 0.7495 - val_recall_2: 0.5
385 - 13s/epoch - 409ms/step
Epoch 2/72
32/32 - 12s - loss: 0.4276 - binary_accuracy: 0.8617 - recall_2: 0.384
6 - val_loss: 0.5601 - val_binary_accuracy: 0.8718 - val_recall_2: 0.3
846 - 12s/epoch - 368ms/step
Epoch 3/72
32/32 - 12s - loss: 0.3241 - binary_accuracy: 0.8988 - recall_2: 0.297
4 - val_loss: 0.3080 - val_binary_accuracy: 0.8895 - val_recall_2: 0.4
872 - 12s/epoch - 367ms/step
Epoch 4/72
32/32 - 12s - loss: 0.2658 - binary_accuracy: 0.9096 - recall_2: 0.338
5 - val_loss: 0.2895 - val_binary_accuracy: 0.8757 - val_recall_2: 0.4
103 - 12s/epoch - 368ms/step
Epoch 5/72
32/32 - 11s - loss: 0.2482 - binary_accuracy: 0.9156 - recall_2: 0.328
2 - val_loss: 0.3552 - val_binary_accuracy: 0.8679 - val_recall_2: 0.6
154 - 11s/epoch - 345ms/step
```



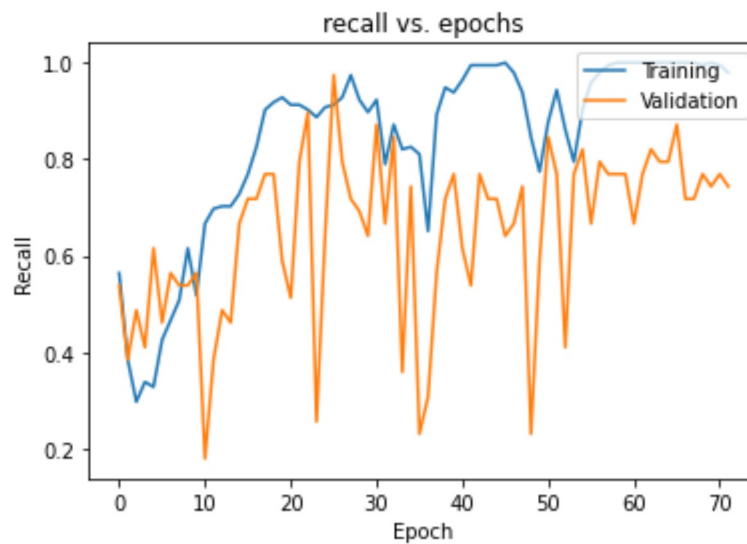
In [26]:

```
plt.plot(list(history.history.values())[0])
plt.plot(list(history.history.values())[3])
plt.title('loss vs. epochs')
plt.ylabel('loss')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')
plt.show()
```




In [27]:

```
plt.plot(list(history.history.values())[2])
plt.plot(list(history.history.values())[5])
plt.title('recall vs. epochs')
plt.ylabel('Recall')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')
plt.show()
```



```
In [28]: ▶ plt.plot(history.history['binary_accuracy'])
plt.plot(history.history['val_binary_accuracy'])
plt.title('Accuracy vs. epochs')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')
plt.show()
```



```
In [29]: ▶ CNN_predictions = model.predict(image_test)

20/20 [=====] - 1s 40ms/step
```

```
In [30]: ▶ NN_predictions = model1.predict(image_test)

20/20 [=====] - 0s 20ms/step
```

```
In [31]: ▶ from sklearn.metrics import confusion_matrix
cnn_cm = confusion_matrix(np.round(CNN_predictions.flatten()).astype(int),
nn_cm = confusion_matrix(np.round(NN_predictions.flatten()).astype(int),
print(cnn_cm)
print(nn_cm)

[[547  16]
 [ 28  43]]
[[502  23]
 [ 73  36]]
```